

EVT in Action

A first simple data set

Let's start by uploading the required packages. The `evir` package is one of the many packages to use EVT on data. Another good one is `evd`. But since `evir` is more pedagogical, we stick to it.

```
library(evir)
```

Let us consider the `danish` dataset, containing large fire insurance claims (expressed in 1000) in Denmark from Thursday 3rd January 1980 until Monday 31st December 1990. The data are available in the `evir` package.


```
data(danish)
head(danish)
```

```
## [1] 1.683748 2.093704 1.732581 1.779754 4.612006 8.725274
```

The first thing to do is to have a look at the data and at some basic statistics.

```
hist(danish, 30, col=2) # Try to produce a better plot with ggplot2
```

Histogram of danish



```
summary(danish)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   1.000   1.321   1.778   3.385   2.967 263.250
```

```
sd(danish)
```


```
## [1] 8.507452
```

As expected (we are considering claims, so only disbursements for us, if we are an insurance company), the data are asymmetric (look at the range and the inter-range) and skewed-to-the-right or positively skewed (the mean is indeed larger than the median). The standard deviation is quite large, compared to the mean, indicating a sensible variability.

With a simple exponential QQ-plot, we can try to understand if heavy tails are present or not. Given the things we have just seen, we would say yes. But let us verify.

The function `qplot` in the `evir` package allows us easily have the plot. The function is built on a GPD, hence the exponential is easily obtained by setting the parameter ξ to 0.


```
qplot(danish,xi=0)
```



The clear concavity in the plot is a strong signal of the presence of heavy tails.

What about a Zipf plot, to look for the behavior of the survival function? The plot is easily made with the function `emplot` (empirical plot). It is important to specify the option `xy` to have a log-log representation.

```
emplot(danish,'xy') # Zipf plot
```

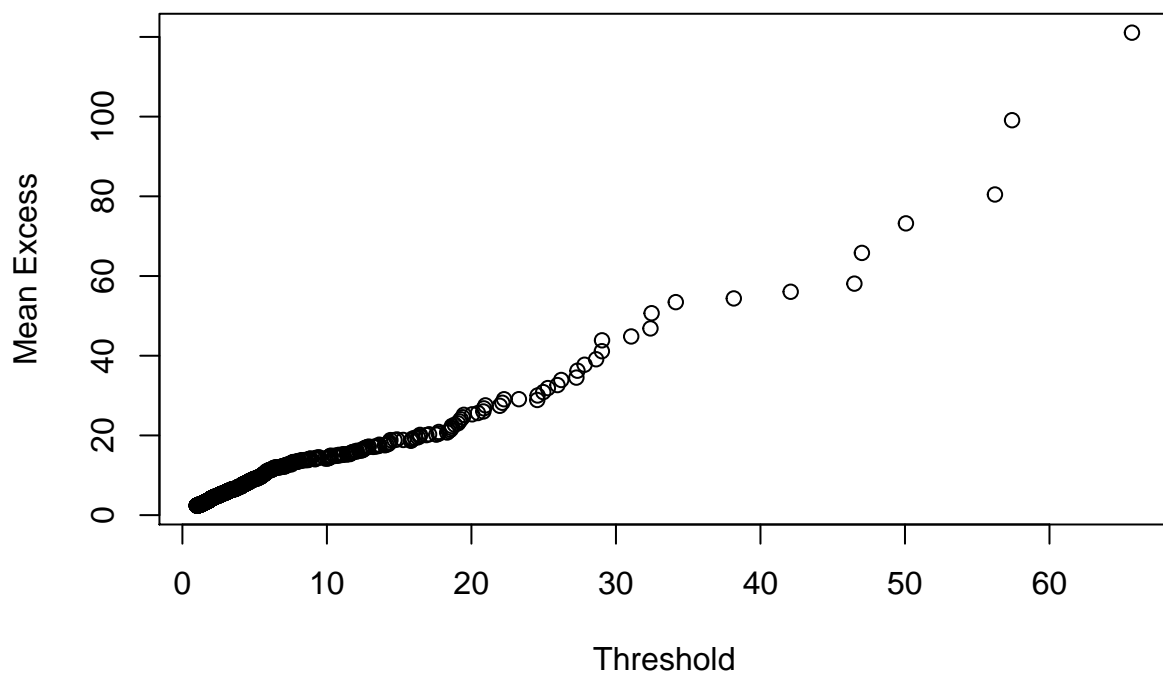


A clear negative linear slope is present. This is a first signal of the fat tailed nature of the data. But remember: a Zipf plot verifies a necessary yet not sufficient condition! Looking at the range of the plot, the credibility of the plot seems ok. The mean and the variance are below 5, while we observe a maximum which is two orders of magnitude larger.

Given that linearity appears from the very beginning, we can think that our Danish claims actually follow a pure power law.

But a Zipf plot is not enough. Let us also consider a Meplot, using the homonymous function `meplot`.


```
meplot(danish)
```



The plot is consistent with van der Wijk's law. Another signal of the presence of a fat tail.

A concentration profile (CP) is another reliable tool to better understand the tail. To build a CP, we can use the functions in the `ineq` library.

```
library(ineq)
sort_danish=sort(danish) # We sort the data
n=length(danish)
CP=c() #Empty array for storage
for (i in 1:n) {
  CP[i]=ineq(sort_danish[i:n],type="Gini") # Truncated Gini
}
plot(1:n,CP,ylim=c(0,1))
```



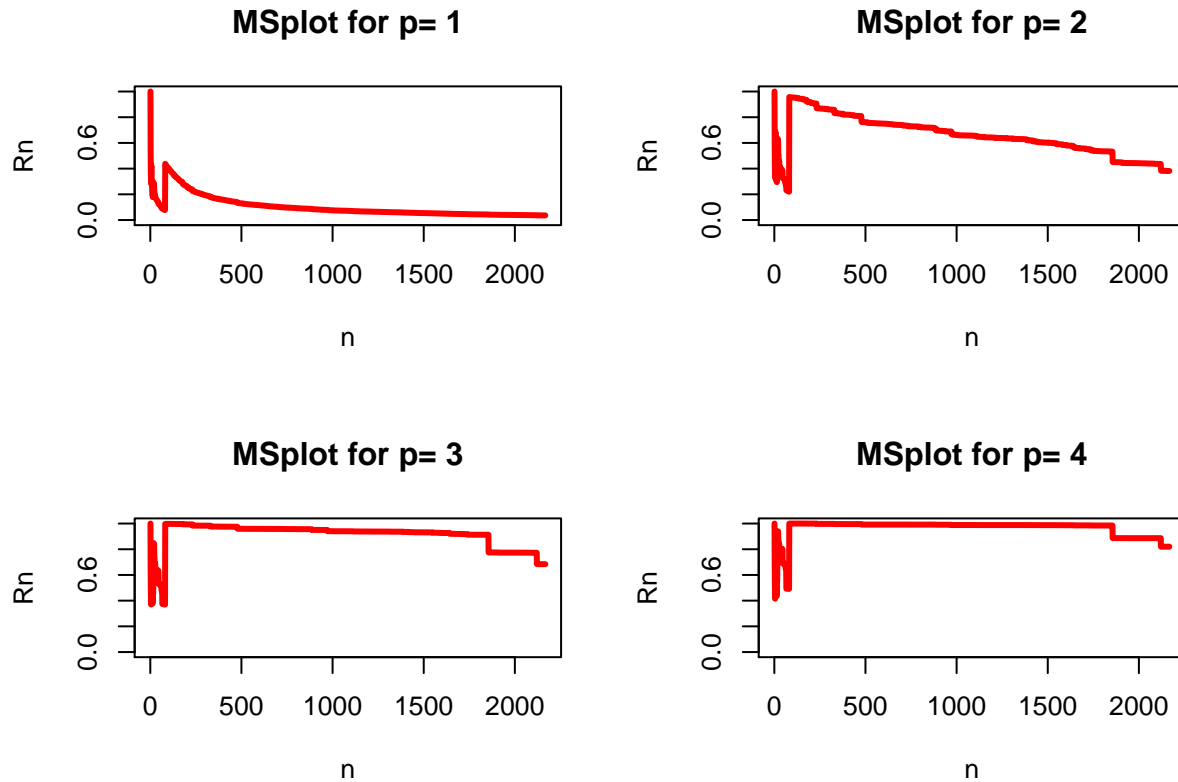
The nearly horizontal behavior we observe (the last part of the plot is to be ignored for the limited amount of points considered) can be seen as a further signal of Paretianity.

Let us try to say something about moments using the MS plot. Here below a simple code to check for the first 4 moments.

```
MSplot <- function(data,p=4) {
  par(mfrow = c(2, 2))
  x=abs(data)
  for (i in 1:p) {
    y=x^i
    S=cumsum(y)
    M=cummax(y)
    R=M/S
    plot(1:length(x),R,type='l', col=2, lwd=3, ylim=c(0,1),xlab='n', ylab='Rn',
        main=paste("MSplot for p=",i))
  }
  par(mfrow = c(1, 1))
  # return(R)
}
```

Let us run it.

MSplot(danish)



We can see that for the first moment convergence is clear, while for the others (starting from the second) we can suspect that they are not defined. If confirmed, a similar finding would tell us that the standard deviation we have computed above is useless for inference. In estimating ξ , we would expect a value between 0.5 and 1.

```
fit=gpd(danish,5)
fit
```

```
## $n
## [1] 2167
##
## $data
## [1] 8.725274 7.898975 7.320644 11.374817 26.214641 14.122076
## [7] 5.424253 11.713031 12.465593 5.875902 5.417277 17.569546
## [13] 7.320644 7.320644 13.620791 21.961933 5.563852 263.250366
## [19] 9.882870 6.319914 5.417277 7.613470 6.674817 19.070278
## [25] 6.325037 5.402635 5.314788 5.856515 19.472914 6.915630
## [31] 8.256881 34.141547 6.888453 9.174312 20.969856 5.242464
## [37] 8.551769 12.895151 8.777628 56.225426 7.863696 7.558322
## [43] 5.111402 5.242464 10.222805 14.678899 5.937759 7.109870
## [49] 7.693316 8.453735 6.422018 50.065531 5.698583 10.178024
## [55] 10.820452 5.469679 24.970273 5.231867 11.890606 5.695600
## [61] 7.074911 20.049941 5.507729 5.001735 65.707491 27.262595
## [67] 15.926278 5.469679 22.258226 5.588585 6.234705 5.561735
## [73] 10.011123 10.072303 7.992070 5.925473 12.631813 6.916554
## [79] 5.300504 5.672970 13.348165 11.431591 11.123471 9.314136
## [85] 5.026178 11.623037 14.293194 13.623037 5.445026 18.646484
## [91] 15.811518 19.162304 18.848168 5.305578 7.235602 7.643979
## [97] 7.539267 5.026178 22.137567 16.300000 6.143355 46.500000
```

```

## [103] 6.200000 7.085000 6.306654 6.563000 10.500000 5.500000
## [109] 9.200000 12.225000 14.239000 5.200000 57.410636 5.850000
## [115] 14.300000 6.167000 13.500000 6.700000 10.700000 19.400000
## [121] 7.230000 8.710000 5.600000 5.207329 6.798457 12.054002
## [127] 5.207329 16.441659 29.026037 12.536162 5.785921 5.400193
## [133] 5.785921 6.075217 18.322083 5.785921 5.593057 8.678881
## [139] 8.100289 10.270010 17.068467 17.743491 5.496625 23.283859
## [145] 6.586271 7.142857 11.131725 12.523191 5.751391 5.565863
## [151] 5.751391 8.812616 32.467532 29.037106 18.552876 16.883117
## [157] 7.792208 27.829314 5.899814 12.059369 6.307978 9.461967
## [163] 7.606679 11.595547 5.194805 5.102041 7.101113 16.415262
## [169] 5.094055 6.015972 18.424135 7.985803 6.511091 38.154392
## [175] 5.210293 5.767524 6.832298 7.542147 5.681455 9.228039
## [181] 27.338066 7.542147 6.140195 11.801242 25.288376 10.204082
## [187] 5.989352 20.452529 5.323869 7.098492 47.019521 24.578527
## [193] 15.882875 7.542147 5.501331 25.953860 8.873114 6.211180
## [199] 10.825200 5.767524 31.055901 5.376799 24.555461 42.091448
## [205] 14.394581 20.863675 9.229467 5.927180 5.664691 14.394581
## [211] 10.137172 6.773920 5.715495 6.435224 5.503810 5.080440
## [217] 5.080440 12.801863 16.088061 152.413209 6.287045 14.013548
## [223] 8.367485 9.398815 5.770533 5.249788 12.701101 11.685013
## [229] 6.011854 32.387807 10.584251 18.628281 7.425743 10.998350
## [235] 8.250825 12.376238 5.198020 5.612211 15.284653 10.184818
## [241] 6.372937 8.085809 6.200495 13.201320 20.826733 5.940594
## [247] 8.250825 14.851485 144.657591 28.630363 19.265677 5.528053
## [253] 5.775578 17.739274
##
## $threshold
## [1] 5
##
## $p.less.thresh
## [1] 0.8827873
##
## $n.exceed
## [1] 254
##
## $method
## [1] "ml"
##
## $par.ests
##          xi          beta
## 0.6320499 3.8074817
##
## $par.ses
##          xi          beta
## 0.1117143 0.4637269
##
## $varcov
##          [,1]      [,2]
## [1,] 0.01248007 -0.03203283
## [2,] -0.03203283 0.21504268
##
## $information
## [1] "observed"

```

```
##
## $converged
## [1] 0
##
## $nllh.final
## [1] 754.1115
##
## attr(,"class")
## [1] "gpd"
# plot(fit) # Commented to produce the pdf.
```


A slightly more complicated case.

Let us now consider the BMW data set in the `evir` package.


```
data(bmw)
? bmw
head(bmw)

## [1] 0.047704097 0.007127223 0.008883307 -0.012440569 -0.003569961
## [6] 0.000000000

plot(bmw)
```




```
plot.ts(bmw)
```



```
hist(bmw, col=4)
```


Histogram of bmw



```
losses=-bmw[bmw<0]
profits=bmw[bmw>=0]
```


```
hist(losses,col=2)
```


Histogram of losses



```
hist(profits,col=3)
```

Histogram of profits



```
summary(losses); sd(losses)
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.    Max.
## 0.0002287 0.0036486 0.0079052 0.0106849 0.0138228 0.1406157
```

```
## [1] 0.01079672
```

```
summary(profits); sd(profits)
```


```
##      Min.   1st Qu.   Median     Mean  3rd Qu.    Max.
## 0.000000 0.001909 0.006211 0.009381 0.013168 0.117192
```

```
## [1] 0.01092088
```


```
summary(bmw)
```

```
##      Min.   1st Qu.   Median     Mean  3rd Qu.    Max.
## -0.1406157 -0.0066560  0.0000000  0.0003407  0.0071261  0.1171918
```


```
qplot(losses,xi=0)
```




```
qplot(profits,xi=0)
```



```
emplot(losses,'xy')
```




```
emplot(profits,'xy')
```




Please notice:

```
emplot(bmw, 'xy')
```


```
## Warning in xy.coords(x, y, xlabel, ylabel, log): 3380 x values <= 0 omitted from
## logarithmic plot
```



```
meplot(losses)
```



```
library(ineq)
sort_losses=sort(losses) # We sort the data
n=length(losses)
CP=c() #Empty array for storage
for (i in 1:n) {
  CP[i]=ineq(sort_losses[i:n],type="Gini") # Truncated Gini
}
plot(1:n,CP,ylim=c(0,1))
```



```


MSplot <- function(data,p=4) {
  par(mfrow = c(2, 2))
  x=abs(data)
  for (i in 1:p) {
    y=x^i
    S=cumsum(y)
    M=cummax(y)
    R=M/S
    plot(1:length(x),R,type='l', col=2, lwd=3, ylim=c(0,1),xlab='n', ylab='Rn',
        main=paste("MSplot for p=",i))
  }
  par(mfrow = c(1, 1))
  # return(R)
}

```


Let us run it.

```
MSplot(losses)
```


MSplot for p= 1




MSplot for p= 2



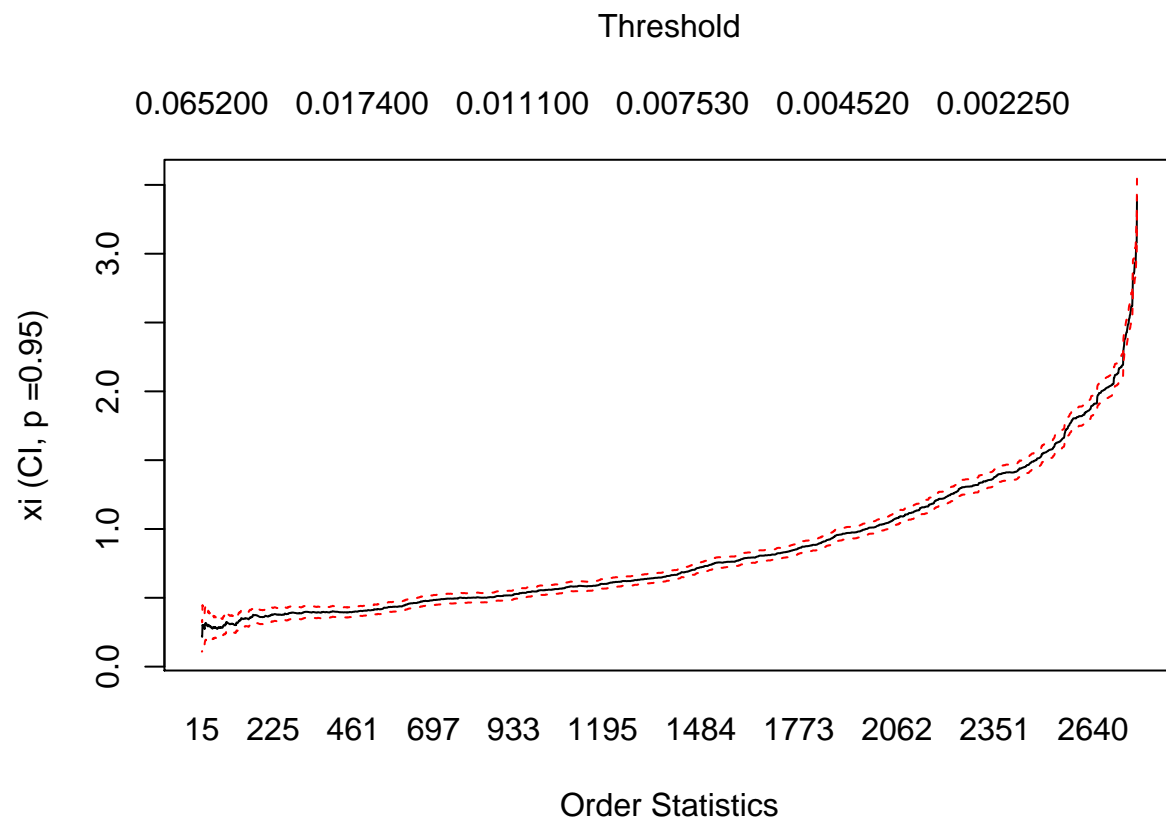
MSplot for p= 3



MSplot for p= 4



```
hill(losses,'xi')
```



```
fit=gpd(losses,0.02)
tail(fit)
```

```
## $par.ests
##          xi          beta
## 0.223297186 0.009248302
##
## $par.ses
##          xi          beta
## 0.0683658225 0.0007703803
##
## $varcov
##           [,1]      [,2]
## [1,] 4.673886e-03 -3.494755e-05
## [2,] -3.494755e-05 5.934858e-07
##
## $information
## [1] "observed"
##
## $converged
## [1] 0
##
## $nllh.final
## [1] -1224.767
```

```
# plot(fit) # Commented to produce the pdf.
```

```
riskmeasures(fit,0.999)
```

```
##           p  quantile      sfall  
## [1,] 0.999 0.1009303 0.1361044
```

The empirical counterparts are

```
quantile(losses,0.9999) #99% Var
```

```
##    99.99%  
## 0.131732
```

```
mean(losses[losses>=quantile(losses,0.9999)]) #99% ES
```

```
## [1] 0.1406157
```

```
fitBM=gev(-bmw,"month")  
fitBM
```

```
## $n.all
```

```
## [1] 6146
```

```
##
```

```
## $n
```

```
## [1] 283
```

```
##
```

```
## $data
```

```
## [1] 0.033019677 0.055259919 0.066891428 0.068877999 0.023835128 0.023092449  
## [7] 0.035115366 0.030063440 0.050842876 0.039941694 0.031013959 0.039396412  
## [13] 0.048453302 0.029417583 0.033783061 0.026425841 0.026598358 0.038631413  
## [19] 0.046887076 0.036826025 0.106175195 0.044687230 0.058654777 0.028176031  
## [25] 0.044666259 0.019492352 0.008697785 0.019301478 0.013453385 0.029693058  
## [31] 0.040710083 0.013700930 0.019466501 0.024850203 0.017798788 0.018079701  
## [37] 0.018323847 0.008458086 0.030940396 0.007513321 0.016826245 0.028967661  
## [43] 0.016654017 0.018659487 0.029137226 0.018070536 0.018235302 0.041225459  
## [49] 0.007976837 0.026597797 0.013237166 0.026068195 0.015605237 0.012444757  
## [55] 0.021876256 0.023924786 0.034657755 0.018418544 0.011214472 0.009001829  
## [61] 0.015595478 0.013506861 0.010298892 0.011157337 0.012895238 0.025355326  
## [67] 0.008765121 0.027107808 0.010206402 0.011275617 0.015205631 0.013417535  
## [73] 0.012360765 0.021851512 0.021111211 0.020871901 0.018310419 0.012981633  
## [79] 0.016119731 0.028002605 0.069554216 0.017148877 0.020496807 0.015062588  
## [85] 0.038596396 0.020995751 0.016289593 0.023379800 0.038451263 0.024264302  
## [91] 0.012085575 0.068078687 0.013428524 0.009963349 0.027324896 0.020475480  
## [97] 0.041947815 0.013293615 0.018199016 0.017977661 0.017773094 0.025637988  
## [103] 0.017038250 0.049572259 0.027796630 0.037448413 0.031355418 0.021699626  
## [109] 0.018211113 0.010931429 0.016123773 0.017726532 0.013882750 0.026525588  
## [115] 0.021081640 0.018482081 0.019868453 0.030210991 0.045415661 0.010078086  
## [121] 0.026954120 0.006155256 0.020924273 0.012877299 0.021631830 0.009804200  
## [127] 0.017655988 0.017050711 0.030511670 0.019076682 0.013192899 0.013831479  
## [133] 0.014225826 0.039844882 0.014622816 0.025401283 0.026486105 0.017457957  
## [139] 0.014043695 0.013402024 0.045258579 0.019724648 0.017905046 0.039381058  
## [145] 0.023808056 0.021783929 0.046091107 0.038087270 0.032365285 0.016007375  
## [151] 0.020097807 0.051884835 0.055613170 0.030047360 0.031916728 0.022432575  
## [157] 0.046816163 0.065191831 0.017443027 0.020191780 0.034982211 0.036689697  
## [163] 0.020569730 0.035401927 0.044916984 0.021291135 0.048802440 0.018587223  
## [169] 0.041909867 0.045651541 0.108521596 0.042561770 0.021587455 0.041483519  
## [175] 0.017679856 0.026193527 0.023709698 0.016034448 0.025052880 0.085731355  
## [181] 0.033531013 0.017630058 0.016876907 0.014403847 0.029542861 0.017305069
```



```
## [187] 0.032251678 0.017839652 0.023564053 0.017936849 0.012596981 0.013623010
## [193] 0.023088421 0.018867592 0.024643272 0.017334021 0.015556229 0.007666020
## [199] 0.022727182 0.021268204 0.013623010 0.020601472 0.019833375 0.140615651
## [205] 0.028951321 0.028348474 0.034301830 0.034748401 0.020399753 0.033601200
## [211] 0.024278141 0.012070965 0.020200330 0.066112864 0.075291383 0.030853653
## [217] 0.043286834 0.034055559 0.014359090 0.022956749 0.022105005 0.018651503
## [223] 0.026770798 0.032213305 0.015825121 0.105774626 0.014809025 0.013005720
## [229] 0.009730427 0.009160196 0.024005906 0.014909519 0.010516251 0.014676834
## [235] 0.027642028 0.020805380 0.043177266 0.029605139 0.058061830 0.047431870
## [241] 0.030041206 0.026409171 0.026883515 0.016880066 0.020540366 0.022792643
## [247] 0.039735321 0.008322718 0.028928575 0.019620427 0.015899478 0.011814395
## [253] 0.024768601 0.022733053 0.034124046 0.015018239 0.022108402 0.014911906
## [259] 0.029923051 0.040018459 0.010374733 0.018794127 0.030126705 0.029413885
## [265] 0.018667209 0.012987196 0.013509372 0.022035888 0.032328887 0.018827493
## [271] 0.022735633 0.018158735 0.019656653 0.020389956 0.031233020 0.034902872
## [277] 0.006851475 0.013087635 0.019297206 0.014371505 0.017753757 0.005153175
## [283] 0.020144566
##
## $block
## [1] "month"
##
## $par.ests
##          xi          sigma          mu
## 0.232488476 0.008931895 0.018679433
##
## $par.ses
##          xi          sigma          mu
## 0.0484898747 0.0004400282 0.0005968189
##
## $varcov
##          [,1]          [,2]          [,3]
## [1,] 2.351268e-03 2.758886e-07 -8.164903e-06
## [2,] 2.758886e-07 1.936248e-07 1.527495e-07
## [3,] -8.164903e-06 1.527495e-07 3.561929e-07
##
## $converged
## [1] 0
##
## $nllh.final
## [1] -850.8664
##
## attr("class")
## [1] "gev"
```

```
# plot(fitBM) # Commented to produce the pdf.
```

Another step:

Let's start by uploading the required packages.

```
library(quantmod)
```

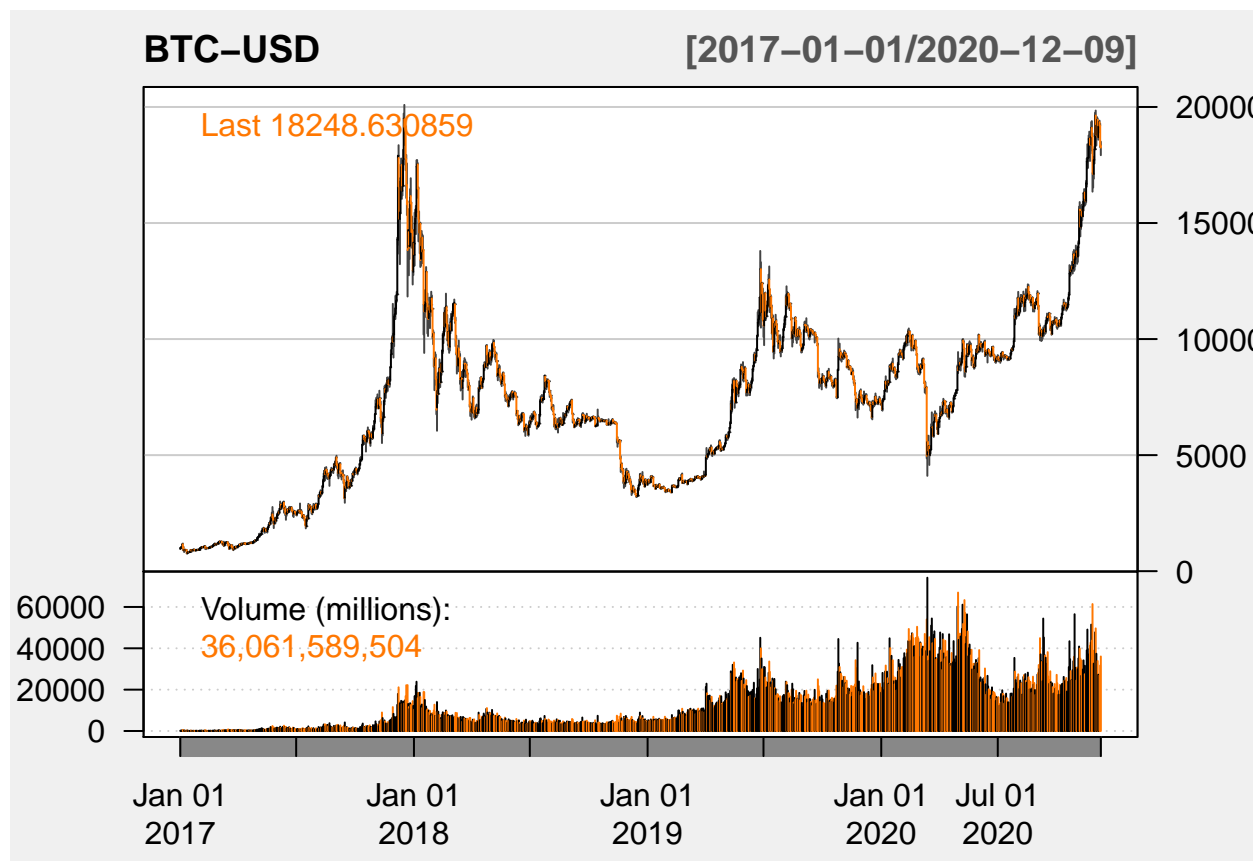
Data exploration

We will study the tails of Bitcoin.

```
getSymbols("BTC-USD", src="yahoo", from="2017-01-01")
```

```
## [1] "BTC-USD"
```

```
chartSeries(`BTC-USD`, up.col="black", theme="white")
```



Let's consider losses and returns.


```
bitcoin=Ad(`BTC-USD`)
lossreturns=diff(log(bitcoin))
losses=-lossreturns[lossreturns<0]
returns=lossreturns[lossreturns>0]
summary(lossreturns)
```

```
##      Index      BTC-USD.Adjusted
##  Min.   :2017-01-01  Min.   :-0.464730
##  1st Qu.:2017-12-26  1st Qu.: -0.014151
##  Median :2018-12-21  Median : 0.002251
##  Mean   :2018-12-21  Mean   : 0.002021
##  3rd Qu.:2019-12-15  3rd Qu.: 0.019602
##  Max.   :2020-12-09  Max.   : 0.225119
##                NA's   :1
```


Some first necessary plots.

```
hist(lossreturns,breaks=20,col=2)
```


Histogram of lossreturns



```
qplot(returns)
```




```
qplot(losses)
```




Zipf Plot and Meplot

```
emplot(returns, 'xy')
```



```
meplot(returns)
```



Continue the analysis for both returns and losses. What do you find? You can share your results/ideas on the course forum. We will comment upon them in the next live class as well.
 #####